




White figure inside each bar is the SPU budget chosen at that α , namely $\min(\text{knee}, \text{banks}/2 \text{ ceiling})$. L40S never reaches its knee — 192 SPU is worth building at every α , so its area and power are flat and only the speedup moves. B200 reaches its knee from 10 ms on and sheds 25 / 63 / 75% of its area at 10 / 20 / 30 ms for under 2% of the runtime. Its speedup is the smaller number only because its GPU baseline is faster: in absolute runtime B200 leads at every α (81.5 ms against 556.2 at 100 μ s).