




★ = knee, the smallest budget within 2% of the floor. Lower is better. At every α and every budget B200 finishes sooner — it shows the smaller speedup only because its GPU baseline is faster too. Curves flatten where the NIC or the SM takes over from the SPU; past that point extra SPUs change nothing. SPU count is not a free parameter; the dashed segments would need a different memory organisation.