




(a) what the SPU-network co-schedule recovers



(b) L40S — 24 ch ceiling (192 SPU)



(c) B200 — 256 ch ceiling (2048 SPU)



(a) Pale: SPU and NIC serialised. Solid: co-scheduled, wall = max(SPU, NIC). Lanes run concurrently, so a pair need not sum to 100%.

(b,c) ★ = knee, the smallest design within 2% of the full card, labelled with its SPU count. L40S has none — 24 channels never reach the network-bound regime. B200 does: at 20 ms, 768 SPU delivers 98% of what 2048 does for 37% of the area. Both gang the communication.