


L40S — 192 SPU

B200 — 2048 SPU



Per-exchange round trip α on the top tick line; the figure below it is the wall-clock gain from the co-schedule.

Pale bar: three lanes serialised. Solid: co-scheduled, wall = max(SPU, SM, NIC). Both gang the communication — same messages, same bytes.

$a^* = \max(\text{SPU}, \text{SM})/(n+2)$ is where the NIC lane overtakes compute. Lanes run concurrently, so the three bars in a group need not sum to 100%.